


# Fast Obstacle Spatial Query on Navigation Mesh

by

Shizhe Zhao



## Minor Thesis

Submitted by Shizhe Zhao

in partial fulfillment of the Requirements for the Degree of  
**Master of Information Technology (Minor Thesis) (3316)**

Supervisor: Assoc. Prof. David Taniar

**Caulfield School of Information Technology**  
**Monash University**

June, 2018

© Copyright

by

Shizhe Zhao

2018

# Contents

<b>Abstract</b> . . . . .	<b>iv</b>
<b>1 Introduction</b> . . . . .	<b>1</b>
1.1 Overview . . . . .	1
1.2 Major Challenges . . . . .	2
1.3 Major Objectives . . . . .	2
1.4 Thesis Organisation . . . . .	2
<b>2 Literature Review</b> . . . . .	<b>3</b>
2.1 Overview . . . . .	3
2.2 Classic pathfinding . . . . .	3
2.3 Spatial Index . . . . .	4
2.3.1 <i>R-tree</i> . . . . .	4
2.3.2 Nearest Neighbor Query . . . . .	5
2.4 Obstacle k-Nearest Neighbor . . . . .	5
2.4.1 In-main-memory OkNN . . . . .	5
2.4.2 Local Visibility Graph . . . . .	5
2.4.3 Fast filter . . . . .	7
2.4.4 Discussion . . . . .	7
2.5 Pathfinding on Navigation Mesh . . . . .	8
2.5.1 Historic background . . . . .	8
2.5.2 Polyanya . . . . .	9
2.6 Other Obstacle Spatial Queries . . . . .	10
2.7 Summary . . . . .	11
<b>3 Proposed Algorithms</b> . . . . .	<b>12</b>
3.1 Overview . . . . .	12
3.2 Interval Heuristic . . . . .	12
3.3 Target Heuristic . . . . .	12
3.4 Summary . . . . .	12
<b>4 Empirical Analysis</b> . . . . .	<b>13</b>
4.1 Overview . . . . .	13
4.2 Benchmark . . . . .	13
4.3 Competitors . . . . .	13
4.4 Experiment 1: lower bounds on performance . . . . .	13
4.5 Experiment 2: computing more nearest neighbor . . . . .	13
4.6 Experiment 3: changing number of targets . . . . .	13
<b>5 Conclusion and Future Work</b> . . . . .	<b>14</b>
5.1 Research Contributions . . . . .	14
5.2 Future Works . . . . .	14

<b>Bibliography . . . . .</b>	<b>15</b>
-------------------------------	-----------

# Fast Obstacle Spatial Query on Navigation Mesh

Shizhe Zhao  
szha414@student.monash.edu  
Monash University, 2018

Supervisor: Assoc. Prof. David Taniar

## Abstract

Obstacle k-Nearest Neighbours problem is the k-Nearest Neighbour problem in a two-dimensional Euclidean plane with obstacles (*OkNN*). Existing and state of the art algorithms for *OkNN* are based on incremental visibility graphs and as such suffer from a well known disadvantage: costly and online visibility checking with quadratic worst-case running times. In this research we develop a new *OkNN* algorithm which avoids these disadvantages by representing the traversable space as a collection of convex polygons; i.e. a Navigation Mesh. We then adapt an recent and optimal navigation mesh algorithm, *Polyanya*, from the single-source single-target setting to the the multi-target case. We also give two new and online heuristics for *OkNN*. In a range of empirical comparisons we show that our approach can be orders of magnitude faster than competing methods that rely on visibility graphs.

*Keywords:* Obstacle Nearest Neighbor, kNN, Navigation Mesh, Spatial Search, Obstacle Distance, Obstacle Navigation

# Chapter 1

## Introduction

### 1.1 Overview

Obstacle k-Nearest Neighbor (OkNN) is a common type of spatial analysis query which can be described as follows: given a set of target points and a collection of polygonal obstacles, all in two dimensions, find the  $k$  closest targets to an a priori unknown query point  $q$ . Such problems appear in a myriad of practical contexts. For example, in an industrial warehouse setting a machine operator may be interested to know the  $k$  closest storage locations where a specific inventory item can be found. OkNN also appears in AI path planning field, for example, in competitive computer games, agent AIs often rely on nearest-neighbour information to make strategic decisions such as during navigation, combat or resource gathering.

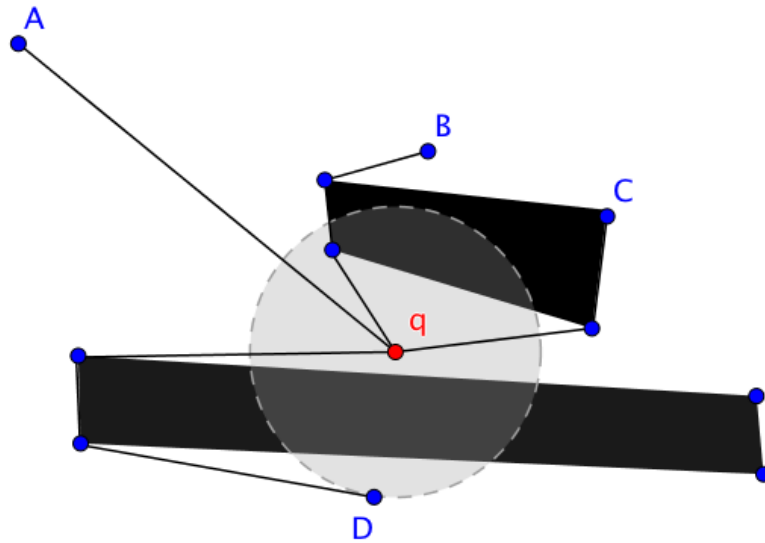


Figure 1.1: We aim to find the nearest neighbour of point  $q$  from among the set of target points  $A, B, C, D$ . Black lines indicate the Euclidean shortest paths from  $q$ . Notice  $D$  is the nearest neighbor of  $q$  under the Euclidean metric but also the furthest neighbor of  $q$  when obstacles are considered.

Traditional kNN queries in the plane (i.e. no obstacles) is a well studied problem that can be handled by algorithms based on spatial index. A key ingredient to the success of these algorithms is the Euclidean metric which provides perfect distance information between any pair of points. When obstacles are introduced however the Euclidean metric becomes an often misleading lower-bound, and the metric becomes obstacle distance. Figure 1.1 shows such an example.

## 1.2 Major Challenges

Two popular algorithms for OkNN, which can deal with obstacles, are *local visibility graphs* [21] and *fast filter* [20]. Though different in details, both of these methods are similar in that they depend on the incremental and online construction of a graph of co-visible points, and use *Dijkstra* to compute shortest path. Algorithms of this type are simple to understand, provide optimality guarantees and the promise of fast performance. Such advantages make incremental visibility graphs attractive to researchers and, despite more than a decade since their introduction, they continue to appear as ingredients in a variety of kNN studies from the literature; e.g. [9–11]. However, incremental visibility graphs also suffer from a number of notable disadvantages including:

1. online visibility checks;
2. an incremental construction process that has up to quadratic space and time complexity for the worst case;
3. duplicated effort, since the graph is discarded each time the query point changes.

In section 2.4, we will introduce these two algorithms with detail, and discuss why they have such disadvantages.

## 1.3 Major Objectives

In this research, we develop a new method for computing *OkNN* which avoids same disadvantages in existing works. Our research extends an existing very fast point-to-point pathfinding algorithm *Polyanya* to multi-target case.

## 1.4 Thesis Organisation

The rest of the thesis is organised as follows:

- In chapter 2, we review related works in different area, includes: AI searching, spatial index and spatial query processing.
- In chapter 3, we introduce the proposed algorithms and discuss their behaviors, formal proof for correctness will be provided.
- In chapter 4, we provide experiment results to demonstrate the performance of proposed algorithms.
- In chapter 5, we summarize our contributions and discuss future works.

## Chapter 2

# Literature Review

### 2.1 Overview

As we've mentioned in previous chapter, OkNN problem appears in both AI path planning and Spatial query processing. Therefore, this literature review includes related works in these two fields.

In section 2.2, we introduce two classic pathfinding algorithms: *Dijkstra* and  $A^*$ , as the historic background. In section 2.3, we introduce a spatial index *R-tree*, and discuss how it solves traditional kNN problem. In section 2.4, we focus on existing works on OkNN, two algorithms based on *Local Visibility Graph* will be discussed. In section 2.5, we introduce a very fast point-to-point algorithm in AI path planning field which shows a new direction to solve OkNN problem. In section 2.6, we briefly discuss other related spatial queries which can be improved by our research.

### 2.2 Classic pathfinding

The most widely used pathfinding algorithm is *Dijkstra* [8]. The algorithm works on a nonnegative weighted graph, it requires a priority queue and regards the length of shortest path as key, and it visit vertices in the order of length of shortest path until requirements be satisfied, e.g. the target has been found. When the target is the furthest vertex to the start vertex, *Dijkstra* has to explore the entire map. Based on such consideration, researchers generalized *Dijkstra* algorithm to *best-first search* which explores a graph by expanding the most promising node chosen according to a specified rule.  $A^*$  [14] is known as a famous *best-first search*, it select the path that minimizes:

$$f(n) = g\text{-value}(n) + h\text{-value}(n)$$

where  $n$  is the last node on the path, *g-value* is the length of shortest path from start to  $n$ , *h-value* is a estimation of shortest path from  $n$  to the goal which is problem-specific. One important property of *h-value* is admissibility, meaning that it never overestimates the actual cost to the target. For example, in an Euclidean plane with obstacles, *h-value* can be the Euclidean distance.

In following sections and the chapter 3, we will show how *Dijkstra* and  $A^*$  algorithms be applied on the OkNN problem.



## 2.3 Spatial Index

### 2.3.1 *R-tree*

*R-tree* has many variations [2, 13, 16, 19], they improve efficiency in different aspects, but they still provide same functionality, so we only introduce the classic *R-tree* in this section.




Figure 2.1: Both segments, circle and irregular shape can be represented by their MBR

*R-tree* is a height-balanced tree [13], all objects are stored in leaf node. In leaf node, if an object is not a point, it would be represented by its *Minimal Bounding Rectangle* (MBR), figure 2.1 shows example of MBRs. Each interior node is also represented by a MBR which contains either leaf nodes or descendant interior nodes. To guarantee efficiency, each non-root node of *R-tree* can contain at least  $m$  entries and at most  $M$  entries, where  $m, M$  are specified constant when *R-tree* is built, and *R-tree*'s root always has two entries. Usually, objects retrieval start from the root, then narrow down to children nodes based on spatial information in their MBRs, and finally retrieve objects from leaf nodes. Figure 2.2 shows how to store and retrieve objects.

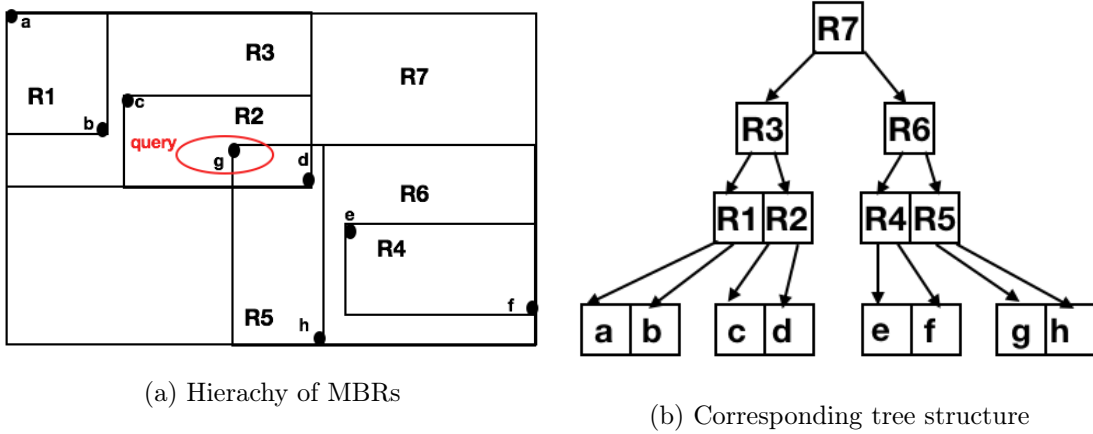


Figure 2.2:  $\{a, b, c, d, e, f, g, h\}$  is the object set,  $R1, R2, R4, R5$  are leaf nodes,  $R3, R6$  are interior nodes, and  $R7$  is the root. The red oval is a range query, starting with  $R7$ , since  $R6$ 's MBR overlapped with query area, we narrow down to  $R6$ , then to  $R5$ , and finally retrieve  $g$ . Notice that  $R3, R2$  also overlap with the query, so they will also be visited, but nothing retrieved.

From the example in figure 2.2, we can see that overlapping area will be explored multiple times in retrieval progress, which duplicated efforts. Thus, some variants use strict non-overlapping interior node (e.g.  $R^+$ -tree [19]), and non-overlapping *R-tree* is a wide topic in spatial index which beyond the scope of this thesis.

### 2.3.2 Nearest Neighbor Query

In the *R-tree*, all nodes are organized by their spatial information, so that the nearest neighbor of a point can be retrieved by exploring tree nodes in some order. To introduce the algorithm, we need to discuss two metrics: given query point  $q$  and the MBR of a tree node

- ***mindist*** is the minimal distance from  $q$  to the MBR, it estimates the distance from  $q$  to inside object, so this metric is the priority of the tree node;
- ***minmaxdist*** is the upper bound of the NN distance of any object inside the MBR, if the *mindist* of any MBR large than this value, then such MBR cannot contains the nearest object, so this metric is for pruning.

The algorithm starts with root node and proceeds down the tree. When it visits a leaf node, current nearest neighbor will be updated; When it visits a non-leaf node, the children of such node is sorted by *mindist*, and pruned by *minmaxdist*, then algorithm does *depth-first-search* on ordered and pruned children nodes; when algorithm finished, the updated nearest neighbor is the answer, and it can be easily generalized to finding kNN by changes below:

- Tracking  $k$  current nearest neighbors, instead of one.
- The pruning is according to the current  $k$ -th nearest neighbor.

This kind of algorithm is called *branch-and-bound* traversal, which has been well studied and widely used in other artificial intelligence areas [19], and existing NN queries are based on it with different ordering and pruning strategies, more details are in [3, 18].

## 2.4 Obstacle k-Nearest Neighbor

In OkNN problem, the metric is obstacle distance, so all existing OkNN algorithms are actually Obstacle Distance Computation (ODC). In following subsections, we introduce these ODC algorithms and show how to apply them on OkNN.

### 2.4.1 In-main-memory OkNN

Solving obstacle path problems in-main-memory has been well studied [6], these works need to compute visibility graph which any pair of co-visible vertices has an edge, figure 2.3 shows an example.

Since all edges of the shortest path belong to visibility graph [17], once precomputed it and include visible edges from query point, we can run shortest path algorithm (e.g. *Dijkstra*) to find  $k$ -nearest neighbor. However, precomputing visibility graph (VG) is costly: even the best algorithm [12] has  $O(m + n \log n)$  runtime, where  $n$  is the number of vertex and  $m$  is the number edges, and in practice  $m$  can reach to  $O(n^2)$ . In spatial database scenario,  $n$  can be more than 10,000, so in-main-memory approach is not suitable for spatial database scenario.

### 2.4.2 Local Visibility Graph

Since building global VG is infeasible, researchers in spatial database field are motivated to design an algorithm that only consider query related area.

In 2004, Zhang proposed the *Local Visibility Graph* (LVG) algorithm [21] to compute obstacle distance. Assume obstacles are stored in *R-tree*, given query point  $q$ , and a target  $t$ , the algorithm runs in following way:






Figure 2.3: The rectangle is the boundary of the map, black polygons are obstacles, and all black lines are edge in the visibility graph.


1. It starts with a small VG centered on  $q$  with radius  $r = d_e(q, t)$  (figure 2.4a);
2. Then compute shortest path on current VG (figure 2.4b);
3. Enlarge the circle to current obstacle distance  $r = d_o(q, t)$ , update the VG incrementally (figure 2.4b) and recompute the shortest path (figure 2.4d);
4. Repeat the previous step until  $r \geq d_o(q, t)$ .




(a) the long rectangle obstacle in the circle is retrieved and included in vg.



(b) current shortest path may be blocked by some obstacles outside



(c) enlarges the circle and updates the vg



(d) recomputes the shortest path

Figure 2.4: LVG algorithm

Since  $r \geq d_o(q, t)$ , when algorithm finish, it guarantees that no obstacle on current shortest path, and thus proof the correctness. This algorithm can be extended to multi-target scenario to solve OkNN problem:

- Initially,  $t$  is the  $k$ -th nearest neighbor in Euclidean space, so that it guarantees that the VG always contains at least  $k$  targets;
- Terminate when  $r$  not less than the current  $k$ -th nearest distance;

### 2.4.3 Fast filter

There's another similar work proposed by Xia [20], the difference is, instead of considering obstacles in a circle area, it only retrieves obstacles on current shortest path, updates the LVG and recomputes the shortest path. Since less obstacles are involved in VG, the algorithm is called *Fast Filter*. Figure 2.5 shows an example.

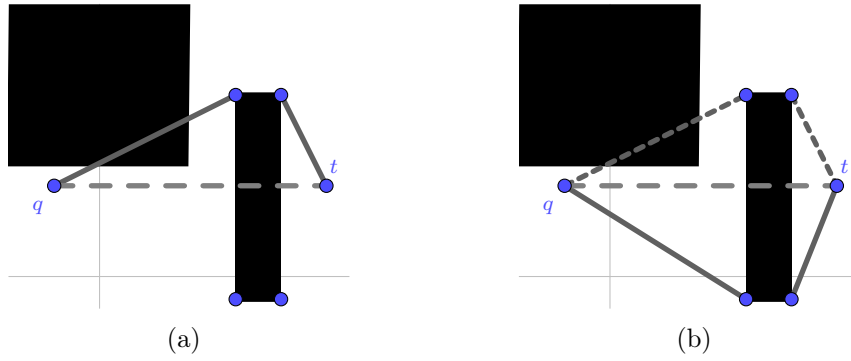


Figure 2.5: Black polygons are obstacle,  $q$  is query point,  $t$  is the target. Initially, the VG only contains the rectangle obstacle between the  $qt$ , and the corresponding shortest path is computed (fig 2.5a). Then retrieves the square obstacle on the current path, updates VG and recomputes shortest path (fig 2.5b).

To solve OkNN, we need following changes:

- initially compute obstacle distance for  $k$  nearest neighbors, and store results in *max-heap* with size  $k$ ;
- keep retrieving next NN and compute the obstacle distance  $d_o$ ;
- when  $d_o$  not large than the top value of heap, pop top and insert the current  $d_o$ ;
- terminate when the Euclidean distance to current NN is large than the Obstacle distance on top of the heap;

### 2.4.4 Discussion

When search space is large (e.g. target is far from the query point), *fast filter* [20] is more efficient because it only consider obstacles that might on the path, meanwhile *LVG* [21] will build visibility graph for a large area, which is slow. In multi-target scenario, when  $k$  is large, *LVG* [21] is more efficient because *dijkstra* is a natural single-source algorithm, meanwhile, each time *fast filter* [20] can only compute path for a single target, so that it duplicates effort for common prefixes of  $k$  targets.

Both *LVG* [21] and *Fast filter* [20] are simple to understand, provide optimality guarantees and the promise of fast performance. Such advantages make them attractive to researchers and, despite more than a decade since their introduction, they continue to appear as ingredients in a variety of kNN studies from the literature; e.g. [9, 10]. However,

these visibility graph based algorithms also suffer from a number of notable disadvantages including:

- (i) costly online visibility checks;
- (ii) an incremental construction process that has up to quadratic space and time complexity for the worst case;
- (iii) duplicated effort, since the graph is discarded each time the query point changes.

## 2.5 Pathfinding on Navigation Mesh

### 2.5.1 Historic background

Because of the limitation of VG, *Navigation Mesh* comes to our sight. Ronald first proposed this concept in 1986 [1], then it is applied on robotics and game pathfinding.

Navigation mesh divides traversable space into convex polygons, the convexity of mesh guarantees that all insides points are co-visible, so that it not needs visibility checking, Figure 2.6 shows an example. Compare to VG (fig 2.3), it can be generated by *Constrained Delaunay Triangulation* [4] in  $O(n \log n)$  times, so that it is feasible to preprocess the entire map. Additionally, adding or removing obstacles only needs to modify a few number of mesh polygons, which is very flexible. It seems like a perfect framework for OkNN, the only problem is that how to compute obstacle distance on it.

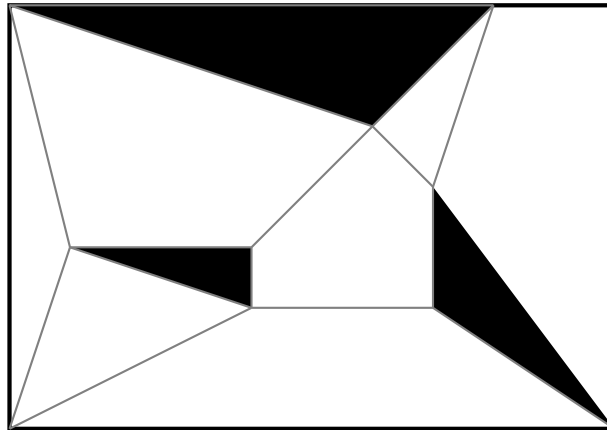


Figure 2.6: Navigation Mesh: The rectangle is the boundary of the map, black polygons are obstacles, gray lines are edges of mesh polygons.

Previous pathfinding algorithms on navigation mesh are not suitable for spatial query processing, and thus it's not attractive to researchers in this fields. Three widely used algorithms are:

- Channel Search [15]: first find an abstract path from start to target composed by polygons, then refine the abstract path to a sequence of concrete points. This algorithm only generate an approximately shortest path, the lack of optimality makes it not suitable for OkNN query.
- TA\* [7]: similar to Channel Search, but it can repeat the search process until finding an optimal shortest path. The repeating is time-consuming, so it is not suitable for query processing.
- TRA\* [7]: similar to TA\*, but TRA\* can utilize preprocessing to speed up pathfinding. The problem is that the precomputed information will be invalid when environment change, so it is not suitable for database scenario.

However, this fact has been changed by a recent work called *Polyanya* [5]. It is a fast, optimal and flexible pathfinding algorithm which is perfect for query processing in spatial database.

### 2.5.2 Polyanya

*Polyanya* can be seen as an instance of  $A^*$ : it performs a best-first search using an admissible heuristic function to prioritise nodes. The mechanical details are however quite different, there are three key components:

- **Search Nodes:** Conventional search algorithms proceed from one traversable point to the next. *Polyanya*, by comparison, searches from one *edge* of the navigation mesh to another. In this model, search nodes are tuples  $(I, r)$  where each  $I = [a, b]$  is a contiguous interval of points and  $r$  is a distinguished point called the *root*. Nodes are constructed such that each point  $p \in I$  is visible from  $r$ . Meanwhile,  $r$  itself corresponds to the last turning point on the path: from  $q$  to any  $p \in I$ . Figure 2.7 shows an example.




Figure 2.7: Search nodes in Polyanya. Notice that the interval  $I = [a, b]$  is a contiguous subset of points drawn from an edge of the navigation mesh. The corresponding root point,  $r$ , is either the query point itself or the vertex of an obstacle. Taken together they form the search node  $(I, r)$ .

- **Successors:** Successor nodes  $(I', r')$  are generated by "pushing" the current interval  $I$  away from its root  $r$  and through the interior of an adjacent and traversable polygon. A successor is said to be *observable* if each point  $p' \in I'$  is visible from  $r$ . The successor node in this case is formed by the tuple  $(I', r)$ . By contrast, a successor is said to be *non-observable* if the *taut* (i.e. locally optimal) path from  $r$  to each  $p' \in I'$  must pass through one of the endpoints of current interval  $I = [a, b]$ . The successor node in this case is formed by the tuple  $(I', r')$  with  $r'$  as one of the points  $a$  or  $b$ . Figure 2.8 shows an example.

Note that the target point is inserted in the open list as a special case (observable or non-observable) successor whenever the search reaches its containing polygon. The interval of this successor contains only the target.

- **Evaluation:** When prioritising nodes for expansion, *Polyanya* makes use of an *f-value* estimation function which bounds the length of the optimal path: from  $q$ , through the current node (i.e. via some  $p \in I$ ) and onto the target. There are three cases to consider which describe the relative positions of the target in relation to the current node. These are illustrated in Figure 2.9. The objective in each case is to choose the unique  $p \in I$  that minimise the estimate. The three cases together are sufficient to guarantee that the estimator is admissible.

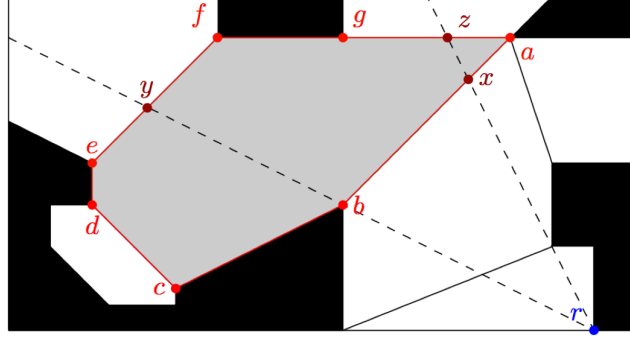


Figure 2.8: From [5]. We expand the node  $([b, x], r)$  which has  $([z, g], r)$  and  $([f, y], r)$  as observable successors. In addition, the nodes  $([c, d], b)$ ,  $([d, e], b)$  and  $([e, y], b)$  are non-observable. All other potential successors can be safely pruned (more details in [5]).

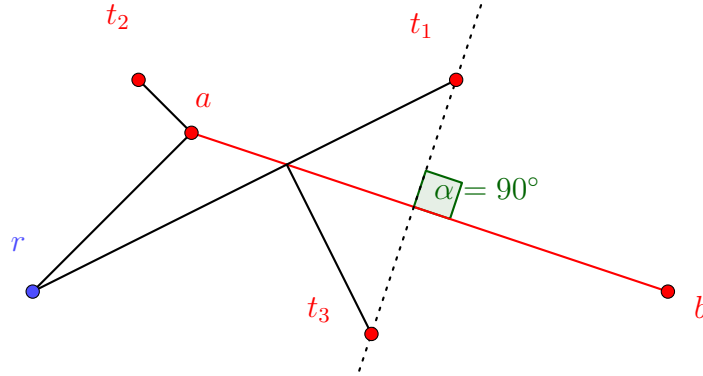


Figure 2.9: Polyanya  $f$ -value estimator. The current node is  $(I, r)$  with  $I = [a, b]$  and each of  $t_1, t_2, t_3$  are possible target locations. **Case 1:** the target is  $t_1$ . In this case the point  $p \in I$  with minimum  $f$ -value is at the intersection of the interval  $I$  and the line  $r \rightarrow t_1$ . **Case 2:** the target is  $t_2$ . In this case the  $p \in I$  with minimum  $f$ -value is one of two endpoints of  $I$ . **Case 3:** the target is  $t_3$ . In this case the  $p \in I$  with minimum  $f$ -value is obtained by first mirroring  $t_3$  through  $[a, b]$  and applying Case 1 or Case 2 to the mirrored point (here,  $t_1$ ). Notice that in this case, simply  $r$  to  $t_3$  doesn't give us the  $h$ -value, based on definition, it must reach the interval first.

Similar to  $A^*$ , *Polyanya* terminates when the target is expanded or when the open list is empty. Extending it to multi-target OkNN is not a trivial problem, we will discuss this in chapter 3.

## 2.6 Other Obstacle Spatial Queries

Obstacle spatial query processing is a broad research area, and many existing works are still based on visibility graph, in this section, we review those works which can get benefit from OkNN in our research.

- **Obstacle Range Query(OR):** given query point  $q$  and range  $r$ , it returns all targets which obstacle distance to  $q$  are less or equal to  $r$  [21]. To solve this by OkNN, we can let  $k$  be infinite and terminate the algorithm when current obstacle distance large than  $r$ .
- **Obstacle Reverse k-Nearest Neighbor(ORkNN):** it is proposed by Gao in 2011 [10] for  $k = 1$ , and be generalized to  $k > 1$  in 2016 [9]. ORkNN given query point  $q$  and  $k$ , return a set of target which regards  $q$  as its OkNN:  $\{t | q \in OkNN(t, k)\}$ . The

query processing has two stages: (i) *search stage*: explore search space to get a set of candidate; (ii) *refine stage*: calls OkNN for each candidate, remove a candidate if  $q$  is not its OkNN. The first stage needs to compute obstacle distance to prune search space, and the second stage can directly get benefit from the improvement of OkNN.

- **Continuous Obstacle k-Nearest Neighbor(COkNN)**: it's proposed in 2009 [11], simliar to OkNN, but the query becomes a segment. The algorithm first generates a list of "split" points, and reduce the problem to finding the OkNN for these points. It also needs to compute obstacle distance to prune search space.

## 2.7 Summary

In this chapter, we review works which related to our research, basically they are in three fields: Spatial Index, Spatial Query Processing and AI Pathfinding. The major finding is that, in OkNN problem, existing VG based algorithms are hard to improve, so we look for a new framework, meanwhile, the new work in AI Pathfinding field shows us a new direction. Such finding inspires us to proposed our research problem: an OkNN algorithm on navigation mesh. We will discuss more detail about this in chapter 3.



## Chapter 3

# Proposed Algorithms

### 3.1 Overview

### 3.2 Interval Heuristic

### 3.3 Target Heuristic

### 3.4 Summary

## Chapter 4

# Empirical Analysis

4.1 Overview

4.2 Benchmark

4.3 Competitors

4.4 Experiment 1: lower bounds on performance

4.5 Experiment 2: computing more nearest neighbor

4.6 Experiment 3: changing number of targets

## Chapter 5

# Conclusion and Future Work

### 5.1 Research Contributions

### 5.2 Future Works

# Bibliography

- [1] Ronald C. Arkin. Path planning for a vision-based autonomous robot. volume 0727, pages 0727 – 0727 – 11, 1987.
- [2] Norbert Beckmann, Hans-Peter Kriegel, Ralf Schneider, and Bernhard Seeger. The r\*-tree: an efficient and robust access method for points and rectangles. In *Acm Sigmod Record*, volume 19, pages 322–331. ACM, 1990.
- [3] King Lum Cheung and Ada Wai-Chee Fu. Enhanced nearest neighbour search on the r-tree. *ACM SIGMOD Record*, 27(3):16–21, 1998.
- [4] L Paul Chew. Constrained delaunay triangulations. *Algorithmica*, 4(1-4):97–108, 1989.
- [5] Michael L Cui, Daniel D Harabor, Alban Grastien, and Canberra Data61. Compromise-free pathfinding on a navigation mesh. *IJCAI*, 2017.
- [6] Mark De Berg, Marc Van Kreveld, Mark Overmars, and Otfried Cheong Schwarzkopf. Computational geometry. In *Computational geometry*, pages 1–17. Springer, 2000.
- [7] Douglas Demyen and Michael Buro. Efficient triangulation-based pathfinding. In *Aaai*, volume 6, pages 942–947, 2006.
- [8] Edsger W Dijkstra. A note on two problems in connexion with graphs. *Numerische mathematik*, 1(1):269–271, 1959.
- [9] Yunjun Gao, Qing Liu, Xiaoye Miao, and Jiacheng Yang. Reverse k-nearest neighbor search in the presence of obstacles. *Information Sciences*, 330:274–292, 2016.
- [10] Yunjun Gao, Jiacheng Yang, Gang Chen, Baihua Zheng, and Chun Chen. On efficient obstructed reverse nearest neighbor query processing. In *Proceedings of the 19th ACM SIGSPATIAL international conference on advances in Geographic Information Systems*, pages 191–200. ACM, 2011.
- [11] Yunjun Gao and Baihua Zheng. Continuous obstructed nearest neighbor queries in spatial databases. In *Proceedings of the 2009 ACM SIGMOD International Conference on Management of data*, pages 577–590. ACM, 2009.
- [12] Subir Kumar Ghosh and David M Mount. An output-sensitive algorithm for computing visibility graphs. *SIAM Journal on Computing*, 20(5):888–910, 1991.
- [13] Antonin Guttman. *R-trees: A dynamic index structure for spatial searching*, volume 14. ACM, 1984.
- [14] Peter E Hart, Nils J Nilsson, and Bertram Raphael. A formal basis for the heuristic determination of minimum cost paths. *IEEE transactions on Systems Science and Cybernetics*, 4(2):100–107, 1968.

- [15] Marcelo Kallmann. Path planning in triangulations. In *Proceedings of the IJCAI workshop on reasoning, representation, and learning in computer games*, pages 49–54, 2005.
- [16] Ibrahim Kamel and Christos Faloutsos. Hilbert r-tree: An improved r-tree using fractals. Technical report, 1993.
- [17] Tomás Lozano-Pérez and Michael A Wesley. An algorithm for planning collision-free paths among polyhedral obstacles. *Communications of the ACM*, 22(10):560–570, 1979.
- [18] Nick Roussopoulos, Stephen Kelley, and Frédéric Vincent. Nearest neighbor queries. In *ACM sigmod record*, volume 24, pages 71–79. ACM, 1995.
- [19] Timos Sellis, Nick Roussopoulos, and Christos Faloutsos. The r+-tree: A dynamic index for multi-dimensional objects. Technical report, 1987.
- [20] Chenyi Xia, David Hsu, and Anthony KH Tung. A fast filter for obstructed nearest neighbor queries. In *British National Conference on Databases*, pages 203–215. Springer, 2004.
- [21] Jun Zhang, Dimitris Papadias, Kyriakos Mouratidis, and Manli Zhu. Spatial queries in the presence of obstacles. *Advances in Database Technology-EDBT 2004*, pages 567–568, 2004.